

Task 1

Use Sqoop tool to export data present in SGOOPOUT folder made while demo of Import table

Table imported with Import command is Person Table

HDFS Files where the table is imported is

Table was imported with 2 MR Jobs with command as below

```
sqoop import --connect jdbc:mysql://localhost/simplidb --table person --
username root -P --target-dir /sqoopoutbyid -m 2
--split-by Person_Id
```

```
[acadmild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root -P --target-dir /sqoopoutByid -m 2 --split-by person_id
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/03/31 05:00:15 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/03/31 05:00:20 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/03/31 05:00:20 INFO tool.CodeGenTool: Beginning code generation
Sat Mar 31 05:00:20 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7
.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServer
Certificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certi
ficate verification.
18/03/31 05:00:22 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/03/31 05:00:22 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/03/31 05:00:22 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadmild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadmild/compile/52b791338ced414f4ba9c6b0382619fd/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/03/31 05:00:27 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadmild/compile/52b791338ced414f4ba9c6b0382619fd/Person.jar
18/03/31 05:00:27 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/03/31 05:00:27 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/03/31 05:00:27 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/03/31 05:00:27 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
18/03/31 05:00:27 INFO mapreduce.ImportJobBase: Beginning import of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadmild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class
]
SLF4J: Found binding in [jar:file:/home/acadmild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/03/31 05:00:28 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/31 05:00:28 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/03/31 05:00:30 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/03/31 05:00:31 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
Sat Mar 31 05:00:35 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7
.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServer
Certificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certi
ficate verification.
18/03/31 05:00:36 INFO db.DBInputFormat: Using read committed transaction isolation
18/03/31 05:00:36 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`person_id`), MAX(`person_id`) FROM `Person`
18/03/31 05:00:36 INFO mapreduce.JobSubmitter: number of splits:2
18/03/31 05:00:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1522393135360_0016
18/03/31 05:00:37 INFO impl.YarnClientImpl: Submitted application application_1522393135360_0016
18/03/31 05:00:37 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1522393135360_0016/
18/03/31 05:00:37 INFO mapreduce.Job: Running job: job_1522393135360_0016
```

```

18/03/31 05:01:27 INFO mapreduce.Job: map 100% reduce 0%
18/03/31 05:01:28 INFO mapreduce.Job: Job job_1522393135360_0016 completed successfully
18/03/31 05:01:28 INFO mapreduce.Job: Counters: 31
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=256014
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=225
    HDFS: Number of bytes written=138
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=4
  Job Counters
    Killed map tasks=1
    Launched map tasks=2
    Other local map tasks=2
    Total time spent by all maps in occupied slots (ms)=53267
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=53267
    Total vcore-milliseconds taken by all map tasks=53267
    Total megabyte-milliseconds taken by all map tasks=54545408
  Map-Reduce Framework
    Map input records=4
    Map output records=4
    Input split bytes=225
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=764
    CPU time spent (ms)=6350
    Physical memory (bytes) snapshot=213876736
    Virtual memory (bytes) snapshot=4126875648
    Total committed heap usage (bytes)=65142784
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=138
18/03/31 05:01:28 INFO mapreduce.ImportJobBase: Transferred 138 bytes in 57.9674 seconds (2.3806 bytes/sec)
18/03/31 05:01:28 INFO mapreduce.ImportJobBase: Retrieved 4 records.

```

```

[acadgild@localhost ~]$ hadoop fs -ls /
18/03/31 05:02:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 24 items
drwxr-xr-x - acadgild supergroup          0 2018-03-30 19:17 /Assignment5_Input
drwxr-xr-x - acadgild supergroup          0 2018-03-30 19:22 /Assignment5_Output1
drwxr-xr-x - acadgild supergroup          0 2018-03-30 20:35 /Assignment5_Output2
drwxr-xr-x - acadgild supergroup          0 2018-03-30 20:51 /Assignment5_Output3
drwxr-xr-x - acadgild supergroup          0 2018-03-31 04:10 /Assignment5_Output4
drwxr-xr-x - acadgild supergroup          0 2018-03-04 22:22 /Prach1
drwxr-xr-x - acadgild supergroup          0 2018-03-24 13:09 /assignment4_input
drwxr-xr-x - acadgild supergroup          0 2018-03-24 14:34 /assignment4_output1
drwxr-xr-x - acadgild supergroup          0 2018-03-24 14:29 /assignment4_output1[A
drwxr-xr-x - acadgild supergroup          0 2018-03-30 12:50 /assignment4_output2
drwxr-xr-x - acadgild supergroup          0 2018-03-30 13:12 /assignment4_output3
drwxr-xr-x - acadgild supergroup          0 2018-03-24 09:38 /hadoopdata
drwxr-xr-x - acadgild supergroup          0 2018-03-30 15:19 /hbase
drwxr-xr-x - acadgild supergroup          0 2018-03-04 19:15 /input
drwxr-xr-x - acadgild supergroup          0 2018-03-05 09:06 /output
drwxr-xr-x - acadgild supergroup          0 2018-03-24 12:22 /output1
drwxr-xr-x - acadgild supergroup          0 2018-03-24 12:32 /output2
drwxr-xr-x - acadgild supergroup          0 2018-03-24 12:37 /output3
drwxr-xr-x - acadgild supergroup          0 2018-03-31 04:54 /sqoopOut_13Apr
drwxr-xr-x - acadgild supergroup          0 2018-03-17 09:59 /sqoopout
drwxr-xr-x - acadgild supergroup          0 2018-02-02 12:49 /sqoopout111
drwxr-xr-x - acadgild supergroup          0 2018-03-31 05:01 /sqoopoutById
drwxrwx-- -- acadgild supergroup          0 2018-03-05 08:40 /tmp
drwxr-xr-x - acadgild supergroup          0 2018-03-18 07:10 /user
[acadgild@localhost ~]$ hadoop fs -ls /sqoopoutById
18/03/31 05:02:17 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r-- 1 acadgild supergroup          0 2018-03-31 05:01 /sqoopoutById/_SUCCESS
-rw-r--r-- 1 acadgild supergroup          59 2018-03-31 05:01 /sqoopoutById/part-m-00000
-rw-r--r-- 1 acadgild supergroup          79 2018-03-31 05:01 /sqoopoutById/part-m-00001
[acadgild@localhost ~]$ hadoop fs -cat /sqoopoutById/part-m-00000
18/03/31 05:02:35 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
1,Shyam,Ram,Patna,Bihar
2,Tanya,Priya,Whitefiled,Bangalore
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /sqoopoutById/part-m-00001
18/03/31 05:02:43 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
3,James,Brown,New York,United States
4,Jhon,Miller, Los Angeles ,United States
[acadgild@localhost ~]$

```

Output is

```

[acadgild@localhost ~]$ hadoop fs -ls /sqoopoutById
18/03/31 05:15:48 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r-- 1 acadgild supergroup          0 2018-03-31 05:01 /sqoopoutById/_SUCCESS
-rw-r--r-- 1 acadgild supergroup          59 2018-03-31 05:01 /sqoopoutById/part-m-00000
-rw-r--r-- 1 acadgild supergroup          79 2018-03-31 05:01 /sqoopoutById/part-m-00001
[acadgild@localhost ~]$

```

```
[acadgild@localhost ~]$ hadoop fs -cat /sqoopoutById/part-m-00000
18/03/31 05:19:04 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
1,Shyam,Ram,Patna,Bihar
2,Tanya,Priya,Whitefield,Bangalore
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /sqoopoutById/part-m-00001
18/03/31 05:19:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
3,James,Brown,New York,United States
4,Jhon,Miller, Los Angeles ,United States
[acadgild@localhost ~]$
```

Now we have export back this to sql.

Solution Approach

- Creating new Database 'Exported' in MySQL
- Export data from HDFS to MySQL

```
Last login: Sat Mar 31 05:15:18 2018 from 192.168.0.3
[acadgild@localhost ~]$ sudo service mysqld start
[sudo] password for acadgild:
Starting mysqld: [ OK ]
[acadgild@localhost ~]$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 63
Server version: 8.0.3-rc-log MySQL Community Server (GPL)

Copyright (c) 2000, 2017, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql>
```

- Created database named as 'Exported'

```
mysql> create database Exported
-> ;
Query OK, 1 row affected (0.73 sec)

mysql> show databases;
+-----+
| Database |
+-----+
| Exported |
| information_schema |
| metastore |
| mysql |
| oozie |
| performance_schema |
| simplidb |
| sys |
+-----+
8 rows in set (0.04 sec)

mysql> 
```

- Set Database 'Exported' as Active Database and list all the tables under that database.

```
mysql> use Exported;
Database changed
mysql> show tables;
Empty set (0.01 sec)

mysql> 
```

- Created new table named as person in database 'Exported'

```
mysql> Create table Person(
-> person_id INT NOT NULL,
-> lname varchar(20),
-> fname varchar(20),
-> area varchar(20),
-> city varchar(20),
-> PRIMARY KEY (person_id)
-> );
Query OK, 0 rows affected (0.70 sec)

mysql> show tables
-> ;
+-----+
| Tables_in_Exported |
+-----+
| Person              |
+-----+
1 row in set (0.01 sec)

mysql> █
```

- No Records in Person Table

```
mysql> select * from Person;
Empty set (0.00 sec)

mysql> █
```

- Command to export the data from HDFS to mysql
scoop export --connect jdbc:mysql://localhost/Exported --table Person --username root -P -export-dir /scoopoutByld

```

[acadgild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost/Exported --table Person --username root -P -export-dir /sqoopoutById
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/03/31 05:29:19 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/03/31 05:29:24 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/03/31 05:29:24 INFO tool.CodeGenTool: Beginning code generation
Sat Mar 31 05:29:24 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7
.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServer
Certificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certi
ficate verification.
18/03/31 05:29:27 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/03/31 05:29:27 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/03/31 05:29:27 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/f404c23cfd7dd69b93909bef16b2a52/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/03/31 05:29:32 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/f404c23cfd7dd69b93909bef16b2a52/Person.jar
18/03/31 05:29:32 INFO mapreduce.ExportJobBase: Beginning export of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class
]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/03/31 05:29:33 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/31 05:29:33 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/03/31 05:29:36 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
18/03/31 05:29:36 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/03/31 05:29:36 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/03/31 05:29:36 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/03/31 05:29:40 INFO input.FileInputFormat: Total input paths to process : 2
18/03/31 05:29:40 INFO input.FileInputFormat: Total input paths to process : 2
18/03/31 05:29:40 INFO mapreduce.JobSubmitter: number of splits:3
18/03/31 05:29:40 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/03/31 05:29:41 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1522393135360_0019
18/03/31 05:29:41 INFO impl.YarnClientImpl: Submitted application application_1522393135360_0019
18/03/31 05:29:42 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1522393135360_0019/
18/03/31 05:29:42 INFO mapreduce.Job: Running job: job_1522393135360_0019
18/03/31 05:30:00 INFO mapreduce.Job: Job job_1522393135360_0019 running in uber mode : false
18/03/31 05:30:00 INFO mapreduce.Job: map_0% reduce_0%
18/03/31 05:30:42 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=383043
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=763
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=18
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Job Counters
    Launched map tasks=3
    Data-local map tasks=3
    Total time spent by all maps in occupied slots (ms)=112326
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=112326
    Total vcore-milliseconds taken by all map tasks=112326
    Total megabyte-milliseconds taken by all map tasks=115021824
  Map-Reduce Framework
    Map input records=4
    Map output records=4
    Input split bytes=512
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=1445
    CPU time spent (ms)=8650
    Physical memory (bytes) snapshot=320831488
    Virtual memory (bytes) snapshot=6183997440
    Total committed heap usage (bytes)=97714176
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=0
18/03/31 05:30:42 INFO mapreduce.ExportJobBase: Transferred 763 bytes in 66.5412 seconds (11.4666 bytes/sec)
18/03/31 05:30:42 INFO mapreduce.ExportJobBase: Exported 4 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

- Validated if records are exported to mysql

```
mysql> select * from Person;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area | city |
+-----+-----+-----+-----+
| 1 | Shyam | Ram | Patna | Bihar |
| 2 | Tanya | Priya | Whitefiled | Bangalore |
| 3 | James | Brown | New York | United States |
| 4 | Jhon | Miller | Los Angeles | United States |
+-----+-----+-----+-----+
4 rows in set (0.00 sec)

mysql>
```

Task2

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table with parameter person_id =3.

Command to Import

```
sqoop import --connect jdbc:mysql://localhost/simplidb --table person --
username root -P --target-dir /sqoopout_where -m 1 --where person_id = 3
```

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root -P --target-dir /sqoopout_where -m 1 --where person_id=3
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/03/31 06:04:32 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/03/31 06:04:37 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/03/31 06:04:37 INFO tool.CodeGenTool: Beginning code generation
Sat Mar 31 06:04:37 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7
.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServer
Certificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certi
ficate verification.
18/03/31 06:04:39 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/03/31 06:04:39 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/03/31 06:04:39 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/432ddd1dbb664505134de68db61d0f3b/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/03/31 06:04:45 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/432ddd1dbb664505134de68db61d0f3b/Person.jar
18/03/31 06:04:45 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/03/31 06:04:45 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/03/31 06:04:45 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/03/31 06:04:45 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
18/03/31 06:04:45 INFO mapreduce.ImportJobBase: Beginning import of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class
]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/03/31 06:04:45 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/31 06:04:45 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/03/31 06:04:47 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/03/31 06:04:48 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
Sat Mar 31 06:04:52 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7
.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServer
Certificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certi
ficate verification.
18/03/31 06:04:53 INFO db.DBInputFormat: Using read committed transaction isolation
18/03/31 06:04:53 INFO mapreduce.JobSubmitter: number of splits:1
18/03/31 06:04:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1522393135360_0020
18/03/31 06:04:54 INFO impl.YarnClientImpl: Submitted application application_1522393135360_0020
18/03/31 06:04:54 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1522393135360_0020/
```



```

18/03/31 06:04:54 INFO mapreduce.Job: Running job: job_1522393135360_0020
18/03/31 06:05:17 INFO mapreduce.Job: Job job_1522393135360_0020 running in uber mode : false
18/03/31 06:05:17 INFO mapreduce.Job: map 0% reduce 0%
18/03/31 06:05:32 INFO mapreduce.Job: map 100% reduce 0%
18/03/31 06:05:33 INFO mapreduce.Job: Job job_1522393135360_0020 completed successfully
18/03/31 06:05:34 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=128157
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=87
    HDFS: Number of bytes written=37
    HDFS: Number of read operations=4
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Other local map tasks=1
    Total time spent by all maps in occupied slots (ms)=12169
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=12169
    Total vcore-milliseconds taken by all map tasks=12169
    Total megabyte-milliseconds taken by all map tasks=12461056
  Map-Reduce Framework
    Map input records=1
    Map output records=1
    Input split bytes=87
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=149
    CPU time spent (ms)=2910
    Physical memory (bytes) snapshot=107864064
    Virtual memory (bytes) snapshot=2063437824
    Total committed heap usage (bytes)=32571392
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=37
18/03/31 06:05:34 INFO mapreduce.ImportJobBase: Transferred 37 bytes in 46.1719 seconds (0.8014 bytes/sec)

```

```

[acadgild@localhost ~]$ hadoop fs -ls /sqoopout where
18/03/31 06:08:37 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup          0 2018-03-31 06:05 /sqoopout where/ SUCCESS
-rw-r--r-- 1 acadgild supergroup        37 2018-03-31 06:05 /sqoopout where/part-m-000000
[acadgild@localhost ~]$ hadoop fs -cat /sqoopout where/part-m-000000
18/03/31 06:08:54 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
3,James,Brown,New York,United States
[acadgild@localhost ~]$ █

```

Created new table Person_copy to export above value

```

mysql> Create table Person_Copy(
-> person_id INT NOT NULL,
-> lname varchar(20),
-> fname varchar(20),
-> area varchar(20),
-> city varchar(20),
-> PRIMARY KEY (person_id)
-> );
Query OK, 0 rows affected (0.25 sec)

mysql> select * from Person_Copy;
ERROR 1146 (42S02): Table 'Exported.Person_Copy' doesn't exist
mysql> select * from Person_Copy;
Empty set (0.03 sec)

mysql> █

```

Export Command –

```

sqoop export --connect jdbc:mysql://localhost/Exported --table Person_Copy --username root -P -
export-dir /sqoopout_where

```



```
[acadmild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost/Exported --table Person_Copy --username root -P -export-dir /sqoopout_where
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6/bin/_hadoop-2.0.4-alpha../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6/bin/_hadoop-2.0.4-alpha../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/03/31 06:11:33 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/03/31 06:11:37 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/03/31 06:11:37 INFO tool.CodeGenTool: Beginning code generation
Sat Mar 31 06:11:38 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7
.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServer
Certificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certi
ficate verification.
18/03/31 06:11:40 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person_Copy` AS t LIMIT 1
18/03/31 06:11:40 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person_Copy` AS t LIMIT 1
18/03/31 06:11:40 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadmild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadmild/compile/1c8c7776d81beb8e9a7f6766e990faa0/Person_Copy.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/03/31 06:11:45 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadmild/compile/1c8c7776d81beb8e9a7f6766e990faa0/Person_Copy.jar
18/03/31 06:11:45 INFO mapreduce.ExportJobBase: Beginning export of Person_Copy
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadmild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class
]
SLF4J: Found binding in [jar:file:/home/acadmild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/03/31 06:11:46 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/31 06:11:46 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/03/31 06:11:48 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
18/03/31 06:11:48 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/03/31 06:11:48 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/03/31 06:11:49 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/03/31 06:11:53 INFO input.FileInputFormat: Total input paths to process : 1
18/03/31 06:11:53 INFO input.FileInputFormat: Total input paths to process : 1
18/03/31 06:11:53 INFO mapreduce.JobSubmitter: number of splits:4
18/03/31 06:11:53 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/03/31 06:11:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1522393135360_0021
18/03/31 06:11:54 INFO impl.YarnClientImpl: Submitted application application_1522393135360_0021
18/03/31 06:11:55 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1522393135360_0021/
18/03/31 06:11:55 INFO mapreduce.Job: Running job: job_1522393135360_0021
18/03/31 06:12:14 INFO mapreduce.Job: Job job_1522393135360_0021 running in uber mode : false
18/03/31 06:12:14 INFO mapreduce.Job: map 0% reduce 0%
```

```
18/03/31 06:13:08 INFO mapreduce.Job: Job job_1522393135360_0021 completed successfully
18/03/31 06:13:09 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=510792
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=700
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=19
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Job Counters
    Launched map tasks=4
    Data-local map tasks=4
    Total time spent by all maps in occupied slots (ms)=200084
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=200084
    Total vcore-milliseconds taken by all map tasks=200084
    Total megabyte-milliseconds taken by all map tasks=204886016
  Map-Reduce Framework
    Map input records=1
    Map output records=1
    Input split bytes=586
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=2263
    CPU time spent (ms)=11770
    Physical memory (bytes) snapshot=429101168
    Virtual memory (bytes) snapshot=8248033280
    Total committed heap usage (bytes)=130285568
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=0
18/03/31 06:13:09 INFO mapreduce.ExportJobBase: Transferred 700 bytes in 80.4235 seconds (8.7039 bytes/sec)
18/03/31 06:13:09 INFO mapreduce.ExportJobBase: Exported 1 records.
You have new mail in /var/spool/mail/acadmild
[acadmild@localhost ~]$
```

```
mysql> select * from Person_Copy;  
Empty set (0.00 sec)
```

```
mysql> select * from Person_Copy;  
+-----+-----+-----+-----+-----+  
| person_id | lname | fname | area      | city          |  
+-----+-----+-----+-----+-----+  
|          3 | James | Brown | New York | United States |  
+-----+-----+-----+-----+-----+  
1 row in set (0.00 sec)
```

```
mysql> █
```

MySQLee has subscriptions for the professional edition license. [Learn more about MySQLee subscriptions](#)